

Research Project Playbook

CSS Final Paper — "Long-term reproduction of expertise in online knowledge commons after generative AI"

Supervisor plan · start → submission · designed for a ~1-month window

0. How to use this document

Open it at the start of every work session. Find "**Where you are now**" below, read the current phase's *Objectives* and *Next actions*, and do the next concrete task. When a phase's *Completion criteria* are all met, move on — not before. Stop and think hard at every **DECISION GATE**.

Three honest framing notes from your supervisor:

1. **A month is tight for an original empirical paper on a tens-of-GB dataset.** This plan has almost no slack. That is why it front-loads a *fail-fast feasibility test* (Phase 5) before you commit real time, and why writing starts in parallel with analysis, not after it.
2. **61% of the marks (Criteria 1-3: Originality, Literature, Design) are decided before you run a single model.** The time budget reflects that. Resist the data-scientist instinct to spend three weeks coding and three days writing. That inversion is the most common way strong students score in the middle.
3. **Two things cause instant rejection regardless of quality: a hallucinated reference and prompt-injection/generic-LLM artefacts.** Track R below exists to make the first impossible. Write in your own voice to avoid the second.

Each phase is tagged [**EXPLORATORY**] (reversible; you're learning, not committing), [**COMMITTING**] (you're locking a decision that's expensive to undo), or [**CHECKPOINT**] (external touchpoint).

Where you are now

Phase	Status
0 Orient to the target	☒ done
1 Map the literature	☒ done
2 Read core papers in full	☐ in progress — full-text reads outstanding
3 Validate the gap	☒ largely done (adversarial review passed, narrowly)

Phase	Status
4 Lock question + operationalization	 RQ drafted; definitions not yet locked
5 Data feasibility spike	 not started — this is your highest-priority action
6 → 13	 ahead

Your immediate front: finish the full-text reads (Phase 2), lock the three operationalization decisions (Phase 4), and run the feasibility spike (Phase 5) — in parallel. Schedule the instructor meeting (Phase 6) now for ~end of Week 1.

Three continuous tracks (run across every phase)

Track R — Reference integrity. Install Zotero today. Every source you cite must be one you have actually opened; capture it in Zotero at the moment you read it. Never paste a citation from memory or from an LLM. Pick one style (APA or Nature) now and keep it consistent. *Rationale: a single fabricated or mis-attributed reference means outright rejection.*

Track C — Reproducibility hygiene. From your first script: a git repo, a fixed folder structure (`/data_raw`, `/data_processed`, `/src`, `/output`, `/paper`), fixed random seeds, and a `README` that grows as you work. Comment code as you write it, not later. *Rationale: 5 random papers per session are executed; if yours is picked and doesn't reproduce, Data & Methods is penalised.*

Track W — Write as you go. Do not save "writing" for the end. Capture your literature synthesis as prose in Phase 2, draft the Methods section during Phases 7–8 while decisions are fresh, and draft Results as figures land in Phase 10. Phase 12 is *consolidation*, not first-drafting. *Rationale: the highest-weighted criteria live in the prose; a paper written in the last 3 days reads like one.*

The phases

Phase 0 — Orient to the target [EXPLORATORY]

- **Purpose:** Know precisely what is being graded before generating anything.
- **Why now:** Everything downstream is optimised against the rubric; you can't aim without the target.
- **Deliverables:** A one-page internalised model of the rubric (weights, tripwires, the lecturer's priorities). (*Done in our discussion.*)
- **Tasks:** Re-read the rubric; note the 61/39 front/back split; note the instant-reject rules and the reproducibility check.
- **Questions to answer:** What separates a top paper from an average one here? What are the hidden expectations?

- **Common mistakes:** Treating this as a data-science project; skimming the rubric once and never returning.
- **Completion criteria:** You can state, from memory, where the points are and what gets you rejected.
- **Rubric contribution:** Indirect but total — it calibrates every later decision.
- **Time:** ½ day (done).

Phase 1 — Map the literature ☒ [EXPLORATORY]

- **Purpose:** Build the terrain map of the field so you can locate an opening.
- **Why now:** You can't judge originality without knowing what exists.
- **Deliverables:** A structured map of sub-areas and the key works in each. (*Done.*)
- **Tasks:** Breadth search across venues (CSCW, ICWSM, CHI, WWW, PNAS Nexus, Nature Hum. Behav., MISQ/ISR, NBER); cluster by theme.
- **Questions:** What are the major research conversations? Where is the field mature vs emerging?
- **Common mistakes:** Reading only abstracts and mistaking familiarity for coverage; anchoring on the first interesting paper.
- **Completion criteria:** You can name the clusters and the 2–3 canonical works in each.
- **Rubric contribution:** Literature Review (4); Originality (6) — you can't claim novelty without the map.
- **Time:** 1–2 days (done).

Phase 2 — Read the core papers in full ● [EXPLORATORY]

- **Purpose:** Move from abstract-level awareness to argument-level command of the ~6 papers your contribution stands on or against.
- **Why now:** Your gap claim and your positioning both depend on knowing *exactly* what the closest work did and did not measure. This is also the defence against hallucinating what a paper says.
- **Deliverables:** A one-paragraph structured note per core paper (question / data / method / finding / **what it did NOT do** / how you differ), all in Zotero.
- **Tasks — read these in full text, not abstracts:**
 - Burtch, Lee & Chen (2024, *Sci Rep*) — **the paper to beat**; note it is cross-sectional composition/exit, window ends Mar 2023, no cohort tracking.
 - Del Rio-Chanona, Laurentsyevea & Wachs (2024, *PNAS Nexus*) — activity by experience stratum; "new-user activity relatively stable."
 - Helic et al. (2025, "Stack Overflow Is Not Dead Yet") — content properties (question difficulty/length), 2-yr window.
 - Xu, Nian & Cabral (2020, *Management Science*, "What Makes Geeks Tick?") — **your methodological precedent** for building individual SO trajectories.

- The reputation-validity paper ("Is Reputation on SO Always a Good Indicator? No!") — the caveat that shapes your outcome measure.
- Lave & Wenger (1991, *Situated Learning*) — **your theoretical spine** (legitimate peripheral participation).
- **Questions:** For each: what is the precise seam between their contribution and mine? Does the full text close my gap anywhere the abstract didn't reveal?
- **Common mistakes:** Reading for summary rather than for the *boundary* of each paper; not reading the theory source (Lave & Wenger) and citing it second-hand.
- **Completion criteria:** You can articulate, in two sentences, exactly how your question differs from Burtch — and the full-text reads confirmed the gap rather than closing it.
- **Rubric contribution:** Literature Review (4) directly; Originality (6) via defensible positioning.
- **Time:** 2-3 days (overlaps Phase 4).

Phase 3 — Validate the gap ☒ / confirm [EXPLORATORY] · DECISION GATE

- **Purpose:** Confirm the opening is real *after* the full-text reads, not just after search snippets.
- **Why now:** Committing to a gap that a full paper already occupies is the classic fatal error; you verify before locking.
- **Deliverables:** A written gap statement: the one measurable question the literature has stepped around (the *conversion function* — conditional on entering and persisting, do newcomers still climb from asker to expert answerer?), with the competitors it's distinct from.
- **Tasks:** Cross-check each full-text note against your question; write the "what's known / what's open" paragraph that will seed your Related Work.
- **Questions:** Does any competitor track longitudinal cohort maturation post-ChatGPT? (Current answer: no.) Is the distinction from Burtch a *difference in kind* (conversion vs inflow), not just degree?
- **Common mistakes:** Confusing "no one used my exact dataset" with "no one asked my question"; declaring a gap from absence-of-evidence without adversarial searching.
- **Completion criteria (GATE):** ☒ **PROCEED** if the gap survives full-text reads and is distinct-in-kind from Burtch. ☐ **REVISE** the angle if a full text turns out to occupy the conversion-function question.
- **Rubric contribution:** Originality (6) — this is where that score is won or lost.
- **Time:** ½-1 day.

Phase 4 — Lock the question and operationalization ☐ [COMMITTING]

- **Purpose:** Convert the refined question into three locked measurement decisions, so the design is unambiguous before data touches it.
- **Why now:** These choices determine whether your central confound is defensible; making them *after* seeing results invites p-hacking accusations and kills the

interpretation score.

- **Deliverables:** A locked one-pager: final RQ, hypotheses (H1-H3), and the three definitions below.
 - **(a) Observation window** — fixed length after account creation (e.g., first 12 months), held constant across cohorts for comparability.
 - **(b) Role-transition definition** — what counts as "becoming an answerer" (e.g., first accepted answer; point at which answering reputation exceeds asking reputation).
 - **(c) Reputation decomposition** — split reputation into *answering-derived* vs *asking-derived*; use answer-side measures (accepted answers, net answer score, tag-specific answer reputation), never raw reputation.
- **Tasks:** Write the hypotheses as falsifiable statements; specify the pre- vs post-ChatGPT cohort cut (Nov 2022); specify the topic-substitutability moderator (proxy by public training-data volume per tag).
- **Questions:** Is each hypothesis falsifiable and tied to a specific variable? Does the design let me separate *composition* (who joins) from *progression* (whether they mature)?
- **Common mistakes:** Vague outcome ("engagement"); using raw reputation as expertise; leaving definitions loose "to decide later."
- **Completion criteria:** A stranger could read your one-pager and compute your variables without asking you anything.
- **Rubric contribution:** Research Design (4) directly; Interpretation (3) protected in advance.
- **Time:** 1-2 days.

Phase 5 — Data feasibility spike [EXPLORATORY] · DECISION GATE (kill-gate #1)

- **Purpose:** Before committing the month, prove the data can actually answer the question. This is the single most important risk-reduction step in the plan.
- **Why now:** If trajectories can't be built, if the maturation signal is too noisy, or if censoring guts the recent cohorts, you must know in Week 1 — while pivoting is cheap — not in Week 3.
- **Deliverables:** A throwaway notebook that, on a *sample* (one or two tags, a few thousand users), (i) assigns users to join-cohorts, (ii) reconstructs a basic asker→answerer trajectory, (iii) plots one pre- vs one post-ChatGPT cohort's maturation curve. Plus a one-paragraph verdict.
- **Tasks:** Pull the sample via the Stack Exchange Data Explorer (fast, SQL) *or* a slice of the Internet Archive dump; build the toy panel; eyeball whether any cohort difference is visible; estimate the censoring problem for 2024-25 cohorts.
- **Questions:** Do the required fields exist and join cleanly (Users, Posts, Votes, Badges)? Is there *any* visible signal, or is variance overwhelming? Can I compute my locked variables on real rows?

- **Common mistakes:** Skipping this and building the full pipeline on faith; treating the spike's result as the finding (it's a feasibility check, not analysis).
- **Completion criteria (GATE):** ✓ **PROCEED** if variables are computable and a signal is at least plausible. 🔄 **PIVOT THE OUTCOME** (e.g., swap role-transition for answering-reputation growth) if one metric is intractable but another works. ➖ **FALL BACK** to an alternative question only if the whole design is unworkable — decide this *now*, with the instructor (Phase 6).
- **Rubric contribution:** De-risks Data & Methods (4) and Research Design (4); prevents a late-stage collapse.
- **Time:** 2-3 days. **Do this essentially now.**

Phase 6 — Instructor meeting #1 [CHECKPOINT]

- **Purpose:** Validate direction, scope, and — above all — the identification strategy with the person grading it, while there's still time to act on feedback.
- **Why now:** After the gap is validated and feasibility is demonstrated, so the meeting is high-value ("here is a concrete, feasible design") not low-value ("here is a vague idea"). Meeting before Phase 5 wastes the meeting; meeting after Week 2 wastes the feedback.
- **Prepare before the meeting:** a one-page brief — RQ + H1-H3, the Burtch distinction stated in two sentences, the cohort + topic-substitutability DiD design, the feasibility-spike plot, and your top-two worries (selection/composition; censoring).
- **Feedback to seek (be specific):** (1) Is the conversion-function framing distinct enough from Burtch to read as original? (2) Is the topic-substitutability DiD a credible identification strategy, or does he want something tighter? (3) Is the causal *language* he'll accept for this observational design — where exactly is his line? (He is a causal-inference specialist and a novelty-hype sceptic; ask directly.) (4) Scope: is one Stack Exchange site enough, or does he expect a cross-site check?
- **Common mistakes:** Arriving with an idea instead of a design; asking "is this okay?" instead of targeted questions; not writing the feedback down and acting on it.
- **Completion criteria:** You leave with a clear go / adjust / reframe signal and a noted list of required changes.
- **Rubric contribution:** Protects Originality (6), Design (4), and Interpretation (3) by catching problems the grader cares about early.
- **Time:** ½ day prep + the meeting. Schedule now.

Phase 7 — Finalize research design & pre-register decisions [COMMITTING]

- **Purpose:** Freeze the methodology and estimation strategy in writing before analysis, incorporating instructor feedback.
- **Why now:** Committing the design *before* looking at outcomes is what lets you interpret honestly and defend against fishing.

- **Deliverables:** A written design section draft (Track W starts the Methods here) covering: cohort construction, the fixed window, outcome variables, the cohort-comparison + substitutability-DiD identification, controls (cohort size, pre-trend), and a pre-analysis note listing your hypotheses and specifications *before* estimation.
- **Tasks:** Specify the estimating equations; define the control/comparison logic; write down robustness checks you commit to running (placebo cohorts, alternative windows, alternative expertise measures).
- **Questions:** Does the design separate composition from progression? Is the identification the strongest available given observational, community-side data? What would falsify each hypothesis?
- **Common mistakes:** Designing loosely and letting the data suggest the hypotheses; no pre-committed robustness plan.
- **Completion criteria:** The Methods draft is complete enough that the analysis is now mechanical execution of a written plan.
- **Rubric contribution:** Research Design (4) directly; Interpretation (3) and Limitations (2) set up.
- **Time:** 1-2 days (overlaps Phase 8 start).

Phase 8 — Build the data pipeline [COMMITTING]

- **Purpose:** Turn raw Stack Exchange data into a clean, reproducible cohort-panel dataset.
- **Why now:** Only after the design is frozen — so you build exactly the dataset the design needs, not a speculative one.
- **Deliverables:** A processed panel (one row per user-period or per user, with cohort, window, and all outcome/control variables), produced by scripted, seeded, commented code under git (Track C).
- **Tasks:** Acquire the full dump (Internet Archive) or scoped tags; parse Users/Posts/Votes/Badges/PostHistory; assign cohorts; compute decomposed reputation and role-transition variables; build the tag-substitutability proxy; write data-provenance notes for the paper.
- **Questions:** Is every transformation logged and reproducible? Are edge cases (deleted posts, merged accounts, bots) handled and documented? Does the processed data match spot-checks against SEDE?
- **Common mistakes:** Manual/unlogged cleaning steps; loading tens of GB naively and stalling (sample or chunk); not saving intermediate processed data.
- **Completion criteria:** A single script (or documented sequence) regenerates the analysis dataset from raw inputs end-to-end.
- **Rubric contribution:** Data & Methods (4) — reproducibility is verified empirically.
- **Time:** 3-4 days.

Phase 9 — Exploratory analysis [EXPLORATORY] · DECISION GATE (pivot point)

- **Purpose:** Understand the data's shape and check that the phenomenon is visible before formal estimation.
- **Why now:** Descriptives reveal whether the signal survives at full scale and expose data problems before they contaminate the main models.
- **Deliverables:** Descriptive figures — cohort sizes over time, maturation curves by cohort, the pre-ChatGPT secular trend, substitutable vs niche tag comparison. Clearly labelled *exploratory*.
- **Tasks:** Plot cohort maturation trajectories; visualise the pre-existing decline you must detrend; sanity-check distributions; look for the substitutability gradient descriptively.
- **Questions:** Is the maturation difference visible descriptively? How severe is the composition shift across cohorts? Is the pre-trend steep enough to threaten identification?
- **Common mistakes:** Sliding from exploration into cherry-picking a specification; presenting exploratory plots as confirmatory results.
- **Completion criteria (GATE):** ☒ **PROCEED** to formal estimation if the design still looks answerable. ☐ **ADJUST** (window, outcome, or comparison) if descriptives reveal a fixable problem. ☐ Escalate to instructor if the phenomenon has vanished at scale.
- **Rubric contribution:** Data & Methods (4); guards Interpretation (3) against surprises.
- **Time:** 2–3 days (overlaps Phase 10).

Phase 10 — Main analysis [COMMITTING]

- **Purpose:** Execute the pre-registered identification strategy and its robustness checks.
- **Why now:** After the design is frozen and the data understood — so estimation is confirmation of a plan, not a search.
- **Deliverables:** The main estimates (cohort comparison + substitutability DiD), the committed robustness checks, and the final figures/tables. Begin drafting Results (Track W) as each lands.
- **Tasks:** Run the specified models; run placebo/alternative-window/alternative-measure robustness; quantify uncertainty; assemble captioned, vectorised (PDF/SVG) figures and numbered tables.
- **Questions:** Do results hold across robustness checks? Is the substitutability differential present (the thing that licenses an AI-substitution reading)? What is the effect size, and is it substantively meaningful, not just significant?
- **Common mistakes:** Running specifications until one "works"; reporting significance without magnitude; unlabelled or raster-only figures.
- **Completion criteria:** Every hypothesis has a clear verdict; robustness is done; figures/tables are final and referenced in draft prose.

- **Rubric contribution:** Data & Methods (4); Interpretation (3).
- **Time:** 4-5 days.

Phase 11 — Interpret findings [COMMITTING]

- **Purpose:** Translate estimates into disciplined social-science claims tied to theory.
- **Why now:** Interpretation must follow, not precede, the results — and it's where this rubric (and this lecturer) are strictest.
- **Deliverables:** A written interpretation linking findings to legitimate peripheral participation and to the Burtch/del Rio-Chanona results, with statistical-vs-substantive and association-vs-causation distinctions made explicit.
- **Tasks:** State what the conversion-function evidence means for expertise reproduction; distinguish what the substitutability DiD does and does not license; write the "consistent with, not proof of" framing for the AI-learning mechanism (you observe the community side, not individual AI use).
- **Questions:** Am I claiming exactly what the design supports — no more? Where does association stop and causation begin in my wording? What's the substantive magnitude?
- **Common mistakes:** Causal verbs ("AI caused newcomers to stop learning") on observational data — an automatic penalty here; overstating generality from one platform.
- **Completion criteria:** Every claim in the draft is backed by a specific result and hedged to the design's limits.
- **Rubric contribution:** Interpretation (3) directly; Originality (6) via the theory linkage.
- **Time:** 2 days (overlaps Phase 12).

Phase 12 — Write the paper (consolidate) [COMMITTING]

- **Purpose:** Assemble a coherent research article from the sections drafted across earlier phases.
- **Why now:** Lit review (Phase 2), Methods (Phase 7), and Results (Phase 10) are already drafted via Track W; this phase joins them into an argument and writes Intro, Abstract, Discussion, and Limitations.
- **Deliverables:** A complete 6,000–8,000-word draft in the required structure, professionally typeset (legible ≥ 11 pt font — not Calibri; 2–2.5 cm margins; ≥ 1.15 spacing), vectorised captioned figures, consistent references from Zotero.
- **Tasks:** Write the Introduction (motivate the phenomenon, preview structure); write Related Work as *synthesis* positioning against Burtch; write the Limitations section (below); write Discussion & Conclusions; write the Abstract last; assemble references.
- **Questions:** Does it read as an argument with a claim, or as a sequence of analyses? Is the Burtch distinction unmistakable to a sceptical reader? Is every section labelled per the required structure?

- **Common mistakes:** Leaving writing to the end (produces a "programming report"); a Related Work that lists instead of synthesises; a generic Limitations section.
- **Completion criteria:** A full draft exists, in structure, within word count, with no placeholder text.
- **Rubric contribution:** All six criteria; especially Literature (4), Design (4), and communication quality throughout.
- **Time:** 5-6 days (much already drafted; this is consolidation + Intro/Discussion/Abstract).

Write the Limitations section specifically around your real threats (this is graded, and generic disclaimers score zero): selection/composition across cohorts (the central threat); reputation-as-imperfect-expertise (mitigated by decomposition); right-censoring of recent cohorts; the pre-existing secular decline; platform-policy shocks (the 2023 GenAI ban and moderator strike); and the causal ceiling — you observe the community-side pipeline, not individual AI use, so the design supports association plus a targeted differential, not proof of mechanism.

Phase 13 — Revise, package, audit, submit [COMMITTING] · [CHECKPOINT optional]

- **Purpose:** Turn a complete draft into a submitted paper that passes every hard check.
- **Why now:** Last, because it operates on a finished draft and a finished analysis.
- **Deliverables:** Final paper (PDF), the zipped code+data+README reproducibility package, and a completed submission checklist.
- **Tasks:**
 - **Reference audit (Track R):** open and verify *every* citation — title, authors, venue, year. One fabricated reference = rejection. Do this deliberately, not casually.
 - **Reproducibility package (Track C):** clean the repo; confirm the pipeline runs end-to-end from raw data on a fresh checkout; write the README (how to reproduce main results); zip code + data (or access instructions).
 - **Self-audit for LLM artefacts:** no "here is a possible...", no placeholders, no generic passages that don't engage your specific content.
 - **Proofread** for colloquialisms, undefined jargon, consistent formatting; confirm word count (6,000–8,000, references excluded); confirm figures are vectorised, numbered, captioned, and referenced in text.
 - **Confirm formal exam registration** — unregistered papers are rejected outright.
 - *(Optional CHECKPOINT):* if time and access allow, a brief instructor touchpoint on results/interpretation before final submission — prepare the key figure and your interpretation, seek feedback on whether your causal language is within bounds.
- **Questions:** Does the code reproduce the reported numbers on a clean run? Is every reference real and correct? Would a sceptical reader see the contribution, not "Burtch again"?

- **Common mistakes:** Trusting references without opening them; a repo that runs on your machine but not from scratch; submitting without registration.
 - **Completion criteria:** Every checklist box ticked; clean-run reproduction confirmed; references verified.
 - **Rubric contribution:** Protects against instant rejection; secures Data & Methods (4) and Limitations (2).
 - **Time:** 2-3 days.
-

Reading strategy (what to read, how deeply, when)

- **Phase 1 (breadth, shallow):** many abstracts + intros across the field, to build the map. Depth: skim. Goal: coverage, not command.
- **Phase 2 (depth, deep):** ~6 core papers in full text — the direct competitors (Burtch, del Rio-Chanona, Helic), the methodological precedent (Xu/Nian/Cabral), the measurement caveat (reputation paper), and the theory source (Lave & Wenger, read from the original). Depth: argument-level; write a boundary-note for each.
- **Phase 7-8 (targeted, deep-on-demand):** method papers for the specific techniques you implement (DiD in observational panels; SO trajectory construction). Read only what you need to implement correctly.
- **Throughout:** whenever you'll cite something, read enough to represent its argument honestly and log it in Zotero. Never cite from an abstract or from memory.

Rule of thumb: **breadth early, depth on the few papers your contribution touches, targeted reading during design/build. Depth is for the papers you position against or borrow from — not for everything.**

When to start data, code, and writing (and why)

- **Data — a small sample now (Phase 5), the full pull after the design freezes (Phase 8).** The early sample is *exploratory feasibility*, not commitment — it exists to fail fast. The full pipeline waits until Phase 7 so you build exactly what the frozen design needs, not a speculative dataset you rework.
- **Code — throwaway spike now (Phase 5); production pipeline in Phase 8.** Don't build the real pipeline on faith before feasibility is shown. Once you do build it, it's under git with seeds from line one (Track C).
- **Writing — now, in fragments (Track W).** Lit-review synthesis in Phase 2, Methods in Phase 7, Results in Phase 10, consolidation in Phase 12. Justification: the front-half criteria (61% of marks) live in prose, and a paper drafted in the final days reads rushed and scores accordingly. Early fragmentary writing also surfaces gaps in your reasoning while you can still fix them.

Decision gates at a glance

- **GATE A (end Phase 3):** Does the gap survive full-text reads and stay distinct-in-kind from Burtch? → proceed / revise angle.
- **GATE B (end Phase 5) — kill-gate:** Can the data compute your variables with a plausible signal? → proceed / pivot outcome / fall back (decide with instructor).
- **GATE C (Phase 6):** Instructor go / adjust / reframe on originality and identification.
- **GATE D (end Phase 9) — pivot point:** Does the phenomenon survive at full scale? → proceed / adjust / escalate.

If you hit a stop signal at B or D, the plan is to *pivot within Stack Exchange* (swap the outcome measure) before abandoning the question — the framing and literature work you've done are reusable.

The roadmap (week-by-week)

Effort is parallelised — reading, writing, and analysis overlap. Durations assume sustained work; adjust to your real hours.

Week	Primary phases	Runs in parallel	Milestone / output	Gate
Week 1	Finish Phase 2 (reads) → Phase 4 (lock definitions) → Phase 5 (feasibility spike) → Phase 6 (instructor mtg)	Track R live; Lit-review notes (Track W)	Locked one-pager; feasibility verdict + plot; instructor go-signal	GATE A, B, C
Week 2	Phase 7 (freeze design) → Phase 8 (build pipeline)	Draft Methods (Track W); Track C repo live	Frozen design + pre-analysis note; clean cohort-panel dataset	—
Week 3	Phase 9 (exploratory) → Phase 10 (main analysis)	Draft Results (Track W)	Descriptive figures; main estimates + robustness; final figures/tables	GATE D
Week 4	Phase 11 (interpret) → Phase 12 (consolidate draft) → Phase 13 (audit + submit)	Reference audit; repro package	Full 6–8k draft → verified references → repro zip → submission	—

Dependency chain (each depends on the last): Rubric ▶ Lit map ▶ **Full-text reads** ▶ Gap validated ▶ **Definitions locked** ▶ **Feasibility proven** ▶ Instructor sign-off ▶ **Design frozen** ▶ Dataset built ▶ Exploration ▶ Main analysis ▶ Interpretation ▶ Draft ▶ Verified + reproducible ▶ Submitted.

The **bold** links are the load-bearing ones — a failure there forces a pivot, which is why the gates cluster around them in Week 1.

Immediate next actions (today)

1. Install Zotero; start logging every source (Track R).
2. Finish the full-text reads of the 6 core papers; write one boundary-note each (Phase 2).
3. Write the locked one-pager: RQ, H1-H3, and the three definitions — observation window, role-transition, reputation decomposition (Phase 4).
4. Start the feasibility spike on a small sample (Phase 5) — this is the highest-value thing you can do this week.
5. Email the instructor to schedule the Week 1 meeting; begin the one-page brief for it.

Return here at the start of each session. Find your row in the roadmap, read that phase's *Next actions*, and do the next task.